

II. Working with data in R (presentation)

Data Science Lab, University of Copenhagen

14 March, 2022

Tidyverse package

The tidyverse is a collection of R packages which, among other things, facilitate data handling and data transformation in R. See <https://www.tidyverse.org/> for details.

We must install and load the R package **tidyverse** before we have access to the functions.

- Install package: One option is to go via the *Tools* menu: *Tools* → *Install packages* → write **tidyverse** in the field called *Packages*. This only has to be done once. Otherwise use the **install.packages** function as shown here:

```
install.packages("tidyverse", repos = "https://mirrors.dotsrc.org/cran/")
```

- Load package: Use the **library** command below (preferred), or go to the *Packages* menu in the bottom right window, find **tidyverse** in the list, and click it. This has to be done in every R-session where you use the package.

```
library(tidyverse)
```

People with SCIENCE PC's (Windows) sometimes have problems with the installation step because R tries to install files to a place, where the user doesn't have permissions to save and edit files. You can try this instead:

- When you start RStudio, right-click the icon and choose *Run as administrator*. Perhaps you can now install packages by clicking *Tools* and *Include Packages* as above.
- If not, then the problem may be that RStudio is trying to install to your science drive (H: or \\a00143.science.domain). If so, try the command **.libPaths()**. If it shows two folders - one at the science drive and one locally one on your computer (C:) - then try the command **install.packages("tidyverse", lib=.libPaths()[2])**.

About the working directory

When working on a project, it is important to know *where you are*. The working directory is the path on your computer that R will *try* to access files from.

There are several helpful commands that help you navigate.

```
# show current working directory (cwd)
getwd()
```

```
# absolute path
setwd("~/Desktop/FromExceltoR/")
```

```
# relative path
```

```
setwd('./Presentations')

# go one step back in the directory
setwd('..')

# show folders in cwd
list.dirs(path = ".", recursive = FALSE)

# set working directory absolute path
setwd("~/Desktop/FromExceltoR/Presentations")
```

Import data

Data from Excel files can be imported via the *Import Dataset* facility. You may get the message that the package **readxl** should be installed. If so, then install it as explained for **tidyverse** above.

- Find *Import Data* in the upper right window in RStudio, and choose *From Excel* in the dropdown menu.
- A new window opens. Browse for the relevant Excel file; then a preview of the dataset is shown. Check that it looks OK, and click *Import*.
- Three things happened: Three lines of code was generated (and executed) in the Console, a new dataset now appears in the Environment window, and the dataset is shown in the top left window. Check again that it looks OK.
- Copy the first two lines of code into your R script (or into an R chunk in your Markdown document), but delete line starting with **View** and write instead the name of the dataset, here **downloads**. Then the first 10 lines of the data set are printed.

```
library(readxl)
downloads <- read_excel("downloads.xlsx")
downloads
```

```
## # A tibble: 147,035 x 6
##   machineName userID size time date month
##   <chr>      <dbl> <dbl> <dbl> <dtm> <chr>
## 1 cs18      146579 2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18      995988 7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18      317649 6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18      748501 13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18      955815 356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18      444174 0 0 1995-04-24 00:00:00 1995-04
## 7 cs18      446911 0 0 1995-04-24 00:00:00 1995-04
## 8 cs18      449552 0 0 1995-04-24 00:00:00 1995-04
## 9 cs18      456142 0 0 1995-04-24 00:00:00 1995-04
## 10 cs18     458942 0 0 1995-04-24 00:00:00 1995-04
## # ... with 147,025 more rows
```

R has stored the data in a so-called *tibble*, a type of data frame. Rows are referred to as *observations* or *data lines*, columns as *variables*. The data rows appear in the order as in the Excel file.

A slight digression: If data are saved in a csv file (comma separated values), possibly generated via an Excel sheet, then data can be read with the `read_csv` function. For example, if the data file is called `mydata.csv` and values are separated with commas, then the command

```
mydata <- read.csv("mydata.csv", sep=",")
```

creates a data frame in R with the data. The data frame is *not* a tibble and some of the commands below would not work for such a data frame.

About the data

The dataset is from Boston University and is about www data transfers from November 1994 to May 1995, see <http://ita.ee.lbl.gov/html/contrib/BU-Web-Client.html>.

- It has 147,035 data lines and 6 variables
 - *size* is the download size in bytes, and *time* is the download time in seconds.
-

Extracting variables, simple summary statistics

Variables can be extracted with the `$`-syntax, and we can use squared brackets to show only the first 40, say, values.

```
time_vector <- downloads$time  
time_vector[1:40]
```

```
## [1] 0.493030 0.325608 0.313704 0.582537 0.259252 0.000000 0.000000 0.000000  
## [9] 0.000000 0.000000 0.000000 0.335502 0.284853 0.000000 0.000000 0.000000  
## [17] 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000  
## [25] 0.285665 0.397111 3.410561 0.267474 0.842364 0.903005 2.784645 2.806157  
## [33] 0.990092 0.477629 0.000000 0.000000 0.000000 0.000000 0.988944 0.000000
```

Summary statistics like mean, standard deviation, median are easily computed for a vector.

Examples of R functions for computing summary statistics: `length`, `mean`, `median`, `sd`, `var`, `sum`, `quantile`, `min`, `max`, `IQR`.

```
length(time_vector)
```

```
## [1] 147035
```

```
mean(time_vector)
```

```
## [1] 0.9539674
```

```
sd(time_vector)
```

```
## [1] 14.22557
```

```
median(time_vector)
```

```
## [1] 0
```

```
min(time_vector)
```

```
## [1] 0
```

Notice that more than half the observations have time equal to zero (median is zero).

Data structures: tibble and data.frame

Before we continue with tidyverse, let's look at some highly used data structures in R. Want to know what data type or structure you have, try the function `class`.

```
# Vectors with characters and numeric values
class(downloads$machineName)
```

```
## [1] "character"
```

```
class(downloads$size)
```

```
## [1] "numeric"
```

```
class(downloads)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

You will need to make structures or convert between these in R. In the example below we make a dataframe and a tibble and convert between these.

```
# Make a dataframe from scratch:
downloads2 <- data.frame(machineName=c("cs18","kermit"), rank=c(1,2))
downloads2
```

```
##  machineName rank
## 1         cs18   1
## 2         kermit  2
```

```
class(downloads2)
```

```
## [1] "data.frame"
```

```
# Convert existing object to a dataframe:
downloads2 <- as.data.frame(downloads)
head(downloads2) # head/top of object
```

```
##  machineName userID  size    time      date  month
## 1         cs18 146579  2464 0.493030 1995-04-24 1995-04
## 2         cs18 995988  7745 0.325608 1995-04-24 1995-04
## 3         cs18 317649  6727 0.313704 1995-04-24 1995-04
## 4         cs18 748501 13049 0.582537 1995-04-24 1995-04
## 5         cs18 955815   356 0.259252 1995-04-24 1995-04
## 6         cs18 444174    0 0.000000 1995-04-24 1995-04
```

```
# Convert existing object to a tibble:
downloads2 <- as_tibble(downloads2)
downloads2
```

```
## # A tibble: 147,035 x 6
##   machineName userID  size  time date      month
##   <chr>      <dbl> <dbl> <dbl> <dtm>      <chr>
## 1 cs18      146579  2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18      995988  7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18      317649  6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18      748501 13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18      955815   356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18      444174    0 0      1995-04-24 00:00:00 1995-04
## 7 cs18      446911    0 0      1995-04-24 00:00:00 1995-04
## 8 cs18      449552    0 0      1995-04-24 00:00:00 1995-04
## 9 cs18      456142    0 0      1995-04-24 00:00:00 1995-04
```

```
## 10 cs18          458942      0 0      1995-04-24 00:00:00 1995-04
## # ... with 147,025 more rows

class(downloads2)

## [1] "tbl_df"      "tbl"        "data.frame"

# Make tibble from scratch:
downloads2 <- tibble(machineName=c("cs18","kermit"), rank=c(1,2))
downloads2

## # A tibble: 2 x 2
##   machineName rank
##   <chr>      <dbl>
## 1 cs18        1
## 2 kermit      2
```

Filtering data (selecting rows): filter

The `filter` function is used to make sub-datasets where only certain datalines (rows) are maintained. It's described with *logical expressions* which datalines should be kept in the dataset.

Say that we only want observations with download time larger than 1000 seconds; there happens to be eight such observations:

```
filter(downloads, time > 1000)

## # A tibble: 8 x 6
##   machineName userID      size time date          month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>      <chr>
## 1 cs18      502807  4055821 1275. 1994-12-02 00:00:00 1994-12
## 2 cs18      16653   2573336 1335. 1994-11-22 00:00:00 1994-11
## 3 cs18      957883   2743516 1151. 1994-11-22 00:00:00 1994-11
## 4 cs18       47910   4720220 1749. 1994-11-22 00:00:00 1994-11
## 5 tweetie   223655    245003 1214. 1995-04-13 00:00:00 1995-04
## 6 kermit    576790  14518894 1380. 1995-04-20 00:00:00 1995-04
## 7 kermit    139654   1079731 1129. 1995-02-23 00:00:00 1995-02
## 8 pluto     337530   8674562 1878. 1995-03-13 00:00:00 1995-03

downloads %>%
  filter(time > 1000)

## # A tibble: 8 x 6
##   machineName userID      size time date          month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>      <chr>
## 1 cs18      502807  4055821 1275. 1994-12-02 00:00:00 1994-12
## 2 cs18      16653   2573336 1335. 1994-11-22 00:00:00 1994-11
## 3 cs18      957883   2743516 1151. 1994-11-22 00:00:00 1994-11
## 4 cs18       47910   4720220 1749. 1994-11-22 00:00:00 1994-11
## 5 tweetie   223655    245003 1214. 1995-04-13 00:00:00 1995-04
## 6 kermit    576790  14518894 1380. 1995-04-20 00:00:00 1995-04
## 7 kermit    139654   1079731 1129. 1995-02-23 00:00:00 1995-02
## 8 pluto     337530   8674562 1878. 1995-03-13 00:00:00 1995-03
```

Or say that only want observations with strictly positive download size:

```
downloads2 <- filter(downloads, size > 0)
downloads2
```

```
## # A tibble: 36,708 x 6
##   machineName userID size time date month
##   <chr>      <dbl> <dbl> <dbl> <dtm> <chr>
## 1 cs18      146579 2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18      995988 7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18      317649 6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18      748501 13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18      955815 356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18      596819 15063 0.336 1995-04-24 00:00:00 1995-04
## 7 cs18      169424 2548 0.285 1995-04-24 00:00:00 1995-04
## 8 cs18      386686 1932 0.286 1995-04-24 00:00:00 1995-04
## 9 cs18      783767 7294 0.397 1995-04-24 00:00:00 1995-04
## 10 cs18     788633 4470 3.41 1995-04-24 00:00:00 1995-04
## # ... with 36,698 more rows
```

Notice that this result is assigned to **downloads2**. It has 36,708 data lines. The original data called **downloads** still exists with 147,035 data lines.

Filtering requires *logical predicates*. These are expressions in terms of columns, which evaluate to either TRUE or FALSE for each row. Logical expressions can be combined with logical operations.

- Comparisons: ==, !=, <, >, <=, >=, %in%, is.na
- Logical operations: ! (not), | (or), & (and). A comma can be used instead of &

Here comes two sub-datasets:

```
# Rows from kermit, and with size greater than 200000 bytes are kept.
filter(downloads2, machineName == "kermit", size > 200000)
```

```
## # A tibble: 98 x 6
##   machineName userID size time date month
##   <chr>      <dbl> <dbl> <dbl> <dtm> <chr>
## 1 kermit    157161 498325 0.629 1995-04-13 00:00:00 1995-04
## 2 kermit    734988 271058 17.3 1995-04-22 00:00:00 1995-04
## 3 kermit    388066 435923 29.2 1995-04-22 00:00:00 1995-04
## 4 kermit     34030 642771 4.80 1995-04-12 00:00:00 1995-04
## 5 kermit    327021 724757 4.98 1995-04-12 00:00:00 1995-04
## 6 kermit     38016 561762 9.75 1995-04-05 00:00:00 1995-04
## 7 kermit    277395 404209 11.3 1995-04-05 00:00:00 1995-04
## 8 kermit    576790 14518894 1380. 1995-04-20 00:00:00 1995-04
## 9 kermit     17623 489473 21.2 1995-04-20 00:00:00 1995-04
## 10 kermit   198041 355963 15.3 1995-04-20 00:00:00 1995-04
## # ... with 88 more rows
```

```
# Rows NOT from kermit, and with size greater than 200000 bytes are kept.
filter(downloads2, machineName != "kermit" & size > 200000)
```

```
## # A tibble: 220 x 6
##   machineName userID size time date month
##   <chr>      <dbl> <dbl> <dbl> <dtm> <chr>
## 1 cs18      204764 2691689 0.834 1995-04-26 00:00:00 1995-04
## 2 cs18      397405 215045 1.10 1994-12-15 00:00:00 1994-12
## 3 cs18      809091 226586 3.92 1994-12-15 00:00:00 1994-12
## 4 cs18      779032 1080472 156. 1994-12-11 00:00:00 1994-12
```

```
## 5 cs18      688294 748705 93.1 1994-12-11 00:00:00 1994-12
## 6 cs18      447740 6360764 863. 1994-12-11 00:00:00 1994-12
## 7 cs18      708452 204918 7.07 1994-12-18 00:00:00 1994-12
## 8 cs18      598668 204918 12.7 1994-12-18 00:00:00 1994-12
## 9 cs18      288167 204918 4.98 1994-12-18 00:00:00 1994-12
## 10 cs18     974956 203714 6.13 1994-12-16 00:00:00 1994-12
## # ... with 210 more rows
```

A helpful function to know which machine names are valid can be:

```
# get unique machineName values in downloads2
distinct(downloads2, machineName)
```

```
## # A tibble: 5 x 1
##   machineName
##   <chr>
## 1 cs18
## 2 piglet
## 3 kermit
## 4 tweetie
## 5 pluto
```

And if you are looking for multiple values for a given variable:

```
downloads2 %>% filter(machineName %in% c("kermit", "pluto"), size > 2000000)
```

```
## # A tibble: 8 x 6
##   machineName userID      size  time date      month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>      <chr>
## 1 kermit    576790 14518894 1380. 1995-04-20 00:00:00 1995-04
## 2 kermit    756949 4418124 439. 1995-04-20 00:00:00 1995-04
## 3 kermit    287308 6935603 88.2 1995-04-24 00:00:00 1995-04
## 4 kermit    928227 9523767 171. 1995-02-08 00:00:00 1995-02
## 5 kermit    128147 2743816 216. 1995-02-23 00:00:00 1995-02
## 6 pluto     867173 4670973 230. 1995-03-14 00:00:00 1995-03
## 7 kermit    456524 2836135 127. 1995-03-31 00:00:00 1995-03
## 8 pluto     337530 8674562 1878. 1995-03-13 00:00:00 1995-03
```

Selecting variables: select

Sometimes, datasets has many variables of which only some are relevant for the analysis. Variables can be selected or skipped with the `select` function.

```
# Without the date variable
select(downloads2, -date)
```

```
## # A tibble: 36,708 x 5
##   machineName userID  size  time month
##   <chr>      <dbl> <dbl> <dbl> <chr>
## 1 cs18      146579 2464 0.493 1995-04
## 2 cs18      995988 7745 0.326 1995-04
## 3 cs18      317649 6727 0.314 1995-04
## 4 cs18      748501 13049 0.583 1995-04
## 5 cs18      955815 356 0.259 1995-04
## 6 cs18      596819 15063 0.336 1995-04
## 7 cs18      169424 2548 0.285 1995-04
```

```
## 8 cs18          386686 1932 0.286 1995-04
## 9 cs18          783767 7294 0.397 1995-04
## 10 cs18         788633 4470 3.41  1995-04
## # ... with 36,698 more rows

# Only include the three mentioned variable names
downloads3 <- select(downloads2, machineName, size, time)
downloads3

## # A tibble: 36,708 x 3
##   machineName size time
##   <chr>      <dbl> <dbl>
## 1 cs18        2464 0.493
## 2 cs18        7745 0.326
## 3 cs18        6727 0.314
## 4 cs18       13049 0.583
## 5 cs18         356 0.259
## 6 cs18       15063 0.336
## 7 cs18        2548 0.285
## 8 cs18        1932 0.286
## 9 cs18        7294 0.397
## 10 cs18       4470 3.41
## # ... with 36,698 more rows
```

Notice that we have made a new dataframe, **downloads3** with only three variables.

Transformations of data

Transformations of existing variables in the data set can be computed and included in the data set with the `mutate` function.

We first compute two new variables, download speed (**speed**) and the logarithm of the download size (**logSize**):

```
downloads3 <- mutate(downloads3, speed = size / time, logSize = log10(size))
downloads3

## # A tibble: 36,708 x 5
##   machineName size time speed logSize
##   <chr>      <dbl> <dbl> <dbl> <dbl>
## 1 cs18        2464 0.493 4998.   3.39
## 2 cs18        7745 0.326 23786.  3.89
## 3 cs18        6727 0.314 21444.  3.83
## 4 cs18       13049 0.583 22400.  4.12
## 5 cs18         356 0.259 1373.   2.55
## 6 cs18       15063 0.336 44897.  4.18
## 7 cs18        2548 0.285 8945.   3.41
## 8 cs18        1932 0.286 6763.   3.29
## 9 cs18        7294 0.397 18368.  3.86
## 10 cs18       4470 3.41 1311.   3.65
## # ... with 36,698 more rows
```

We then make a new categorical variable, **slow**, which is “Yes” if speed < 150 and “No” otherwise

```
downloads3 <- mutate(downloads3, slow = ifelse(speed < 150, "Yes", "No"))
downloads3
```



```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18      2464 0.493 4998. 3.39 No
## 2 cs18      7745 0.326 23786. 3.89 No
## 3 cs18      6727 0.314 21444. 3.83 No
## 4 cs18     13049 0.583 22400. 4.12 No
## 5 cs18       356 0.259 1373. 2.55 No
## 6 cs18     15063 0.336 44897. 4.18 No
## 7 cs18      2548 0.285 8945. 3.41 No
## 8 cs18      1932 0.286 6763. 3.29 No
## 9 cs18      7294 0.397 18368. 3.86 No
## 10 cs18     4470 3.41 1311. 3.65 No
## # ... with 36,698 more rows
```

Counting, tabulation of categorical variables: count

The `count` function is useful for counting data datalines, possibly according to certain criteria or for the different levels of categorical values.

```
# Total number of observations in the current dataset
count(downloads3)
```

```
## # A tibble: 1 x 1
##       n
##   <int>
## 1 36708
```

```
# Number of observations from each machine
count(downloads3, machineName)
```

```
## # A tibble: 5 x 2
##   machineName n
##   <chr>      <int>
## 1 cs18      3814
## 2 kermit    9094
## 3 piglet   11200
## 4 pluto     5253
## 5 tweetie   7347
```

```
# Number of observations which have/have not size larger than 5000
count(downloads3, size>5000)
```

```
## # A tibble: 2 x 2
##   `size > 5000` n
##   <lgl>        <int>
## 1 FALSE      25865
## 2 TRUE       10843
```

```
# Number of observations for each combination of machine name and the *slow* variable.
count(downloads3, machineName, slow)
```

```
## # A tibble: 10 x 3
##   machineName slow      n
##   <chr>        <chr> <int>
## 1 cs18        No     3662
```

```
## 2 cs18      Yes      152
## 3 kermit    No       8717
## 4 kermit    Yes       377
## 5 piglet    No      10734
## 6 piglet    Yes       466
## 7 pluto     No      4963
## 8 pluto     Yes       290
## 9 tweetie   No      6983
## 10 tweetie  Yes       364
```

Sorting data: `arrange`

The `arrange` function can be used to sort the data according to one or more columns.

Let's sort the data according to download size (ascending order). The first lines of the sorted data set is printed on-screen, but the dataset **downloads3** has *not* been changed.

```
arrange(downloads3, size)
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl>   <dbl> <chr>
## 1 cs18          3  3.73 0.804   0.477 Yes
## 2 piglet        3  1.53 1.96    0.477 Yes
## 3 piglet        3  1.53 1.96    0.477 Yes
## 4 tweetie       3  1.11 2.71    0.477 Yes
## 5 kermit        3  1.12 2.69    0.477 Yes
## 6 pluto         3  8.60 0.349   0.477 Yes
## 7 pluto         3  9.87 0.304   0.477 Yes
## 8 pluto         3  3.78 0.793   0.477 Yes
## 9 pluto         3  4.68 0.641   0.477 Yes
## 10 pluto        3  4.93 0.608   0.477 Yes
## # ... with 36,698 more rows
```

Two different examples:

```
# According to download size in descending order
arrange(downloads3, desc(size))
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl>   <dbl> <chr>
## 1 kermit    14518894 1380.  10522.   7.16 No
## 2 piglet    14158123  123.  115169.   7.15 No
## 3 kermit     9523767  171.   55562.   6.98 No
## 4 piglet     9384067   80.0 117309.   6.97 No
## 5 pluto     8674562 1878.   4619.   6.94 No
## 6 kermit     6935603   88.2  78655.   6.84 No
## 7 cs18      6360764   863.   7374.   6.80 No
## 8 piglet     5143062   597.   8611.   6.71 No
## 9 piglet     4812334   215.  22345.   6.68 No
## 10 cs18      4720220 1749.   2700.   6.67 No
## # ... with 36,698 more rows
```

```
# After machine name and then according to download size in descending order
arrange(downloads3, machineName, desc(size))
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr> <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18 6360764 863. 7374. 6.80 No
## 2 cs18 4720220 1749. 2700. 6.67 No
## 3 cs18 4055821 1275. 3180. 6.61 No
## 4 cs18 3047343 20.9 146038. 6.48 No
## 5 cs18 2952381 318. 9289. 6.47 No
## 6 cs18 2743516 1151. 2383. 6.44 No
## 7 cs18 2691689 0.834 3228695. 6.43 No
## 8 cs18 2613025 18.5 140959. 6.42 No
## 9 cs18 2573336 1335. 1928. 6.41 No
## 10 cs18 1931453 186. 10388. 6.29 No
## # ... with 36,698 more rows
```

Grouping: group_by

We can group the dataset by one or more categorical variables with `group_by`. The dataset is not changed as such, but - as we will see - grouping can be useful for computation of summary statistics and graphics.

Here we group after machine name (first) *and* the slow variable (second). The only way we can see it at this point is in the second line in the output (`# Groups:`):

```
# Group according to machine
group_by(downloads3, machineName)
```

```
## # A tibble: 36,708 x 6
## # Groups:   machineName [5]
##   machineName size time speed logSize slow
##   <chr> <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18 2464 0.493 4998. 3.39 No
## 2 cs18 7745 0.326 23786. 3.89 No
## 3 cs18 6727 0.314 21444. 3.83 No
## 4 cs18 13049 0.583 22400. 4.12 No
## 5 cs18 356 0.259 1373. 2.55 No
## 6 cs18 15063 0.336 44897. 4.18 No
## 7 cs18 2548 0.285 8945. 3.41 No
## 8 cs18 1932 0.286 6763. 3.29 No
## 9 cs18 7294 0.397 18368. 3.86 No
## 10 cs18 4470 3.41 1311. 3.65 No
## # ... with 36,698 more rows
```

```
# Group according to machine and slow
group_by(downloads3, machineName, slow)
```

```
## # A tibble: 36,708 x 6
## # Groups:   machineName, slow [10]
##   machineName size time speed logSize slow
##   <chr> <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18 2464 0.493 4998. 3.39 No
## 2 cs18 7745 0.326 23786. 3.89 No
## 3 cs18 6727 0.314 21444. 3.83 No
## 4 cs18 13049 0.583 22400. 4.12 No
```

```
## 5 cs18          356 0.259 1373.    2.55 No
## 6 cs18        15063 0.336 44897.    4.18 No
## 7 cs18         2548 0.285  8945.    3.41 No
## 8 cs18         1932 0.286  6763.    3.29 No
## 9 cs18         7294 0.397 18368.    3.86 No
## 10 cs18        4470 3.41  1311.    3.65 No
## # ... with 36,698 more rows
```

Summary statistics, revisited: `summarize`

Recall how we could compute summary statistics for a single variable in a dataset, e.g.

```
mean(downloads3$size)
```

```
## [1] 16638.36
```

```
max(downloads3$size)
```

```
## [1] 14518894
```

With `summarize` we can compute summary statistics for a variable for each level of a grouping variable or for each combination of several grouping variables.

First, a bunch of summaries for the size variable for each machine name, where we give explicit names for the new variables:

```
downloads.grp1 <- group_by(downloads3, machineName)
summarize(downloads.grp1,
  avg = mean(size),
  med = median(size),
  stdev = sd(size),
  total = sum(size),
  n = n())
```

```
## # A tibble: 5 x 6
##   machineName   avg   med  stdev   total     n
##   <chr>       <dbl> <dbl>  <dbl>   <dbl> <int>
## 1 cs18       26375. 1990. 208915. 100593281 3814
## 2 kermit     19247. 2466 213985. 175032552 9094
## 3 piglet     14121. 2146. 188340. 158149841 11200
## 4 pluto      13822. 2069 144425. 72605544 5253
## 5 tweetie    14207. 2197  94318. 104379794 7347
```

Second, the same thing but for each combination of machine name and the slow variable:

```
downloads.grp2 <- group_by(downloads3, machineName, slow)
summarize(downloads.grp2,
  avg = mean(size),
  med = median(size),
  stdev = sd(size),
  total = sum(size),
  n = n())
```

```
## # A tibble: 10 x 7
## # Groups:   machineName [5]
##   machineName slow   avg   med  stdev   total     n
##   <chr>       <chr>  <dbl> <dbl>  <dbl>   <dbl> <int>
```

```
## 1 cs18      No      27445. 2092. 213140. 100503042 3662
## 2 cs18      Yes       594.  368.    614.    90239   152
## 3 kermit    No      20030. 2598  218529. 174602282 8717
## 4 kermit    Yes      1141.  541   3049.    430270   377
## 5 piglet    No      14687. 2264  192365. 157650747 10734
## 6 piglet    Yes      1071.  416.   1934.    499094   466
## 7 pluto     No      14564. 2164  148551.  72280790 4963
## 8 pluto     Yes      1120.  413   2108.    324754   290
## 9 tweetie   No      14894. 2373  96694.  104001733 6983
## 10 tweetie  Yes      1039.  471   2603.    378061   364
```

Third, mean and standard deviation for several variables:

```
summarize_at(downloads.grp2, c("time", "size"), list(ave=mean, stdev=sd))
```

```
## # A tibble: 10 x 6
## # Groups:   machineName [5]
##   machineName slow  time_ave size_ave time_stdev size_stdev
##   <chr>         <chr>   <dbl>   <dbl>    <dbl>    <dbl>
## 1 cs18         No      5.17   27445.    57.1   213140.
## 2 cs18         Yes     9.63    594.    17.8     614.
## 3 kermit       No      3.41  20030.    25.3   218529.
## 4 kermit       Yes    20.7   1141.    47.8   3049.
## 5 piglet       No      2.33  14687.    13.8  192365.
## 6 piglet       Yes    19.4   1071.    40.2   1934.
## 7 pluto        No      3.40  14564.    30.4  148551.
## 8 pluto        Yes    21.7   1120.    46.3   2108.
## 9 tweetie      No      2.68  14894.    17.3   96694.
## 10 tweetie     Yes    17.8   1039.    34.5   2603.
```

The datasets with summaries can be saved as datasets themselves, for example to be used as the basis for certain graphs.

The pipe operator: %>%

Two or more function calls can be evaluated sequentially using the so-called pipe operator, %>%. Nesting of function calls becomes more readable, and intermediate assignments are avoided.

Let's try it to do a bunch of things in one go, starting with the original dataset:

```
downloads %>%
  filter(size>0) %>% # Subset of data
  group_by(machineName) %>% # Grouping
  summarize(ave = mean(size)) %>% # Compute mean
  arrange(ave) # Sort after mean
```

```
## # A tibble: 5 x 2
##   machineName  ave
##   <chr>      <dbl>
## 1 pluto     13822.
## 2 piglet    14121.
## 3 tweetie   14207.
## 4 kermit    19247.
## 5 cs18     26375.
```

More functions from tidyverse

Below are three useful functions for column manipulation, `relocate`, `rename` and `pull`:

```
# relocate (move one or more columns):
```

```
downloads %>% relocate(time, .before = size)
```

```
## # A tibble: 147,035 x 6
```

```
##   machineName userID  time  size date          month
##   <chr>         <dbl> <dbl> <dbl> <dtm>         <chr>
## 1 cs18         146579 0.493 2464 1995-04-24 00:00:00 1995-04
## 2 cs18         995988 0.326 7745 1995-04-24 00:00:00 1995-04
## 3 cs18         317649 0.314 6727 1995-04-24 00:00:00 1995-04
## 4 cs18         748501 0.583 13049 1995-04-24 00:00:00 1995-04
## 5 cs18         955815 0.259 356 1995-04-24 00:00:00 1995-04
## 6 cs18         444174 0      0 1995-04-24 00:00:00 1995-04
## 7 cs18         446911 0      0 1995-04-24 00:00:00 1995-04
## 8 cs18         449552 0      0 1995-04-24 00:00:00 1995-04
## 9 cs18         456142 0      0 1995-04-24 00:00:00 1995-04
## 10 cs18        458942 0      0 1995-04-24 00:00:00 1995-04
## # ... with 147,025 more rows
```

```
# rename (rename one column):
```

```
downloads %>% rename(year.month=month)
```

```
## # A tibble: 147,035 x 6
```

```
##   machineName userID  size  time date          year.month
##   <chr>         <dbl> <dbl> <dbl> <dtm>         <chr>
## 1 cs18         146579 2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18         995988 7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18         317649 6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18         748501 13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18         955815 356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18         444174 0 0 1995-04-24 00:00:00 1995-04
## 7 cs18         446911 0 0 1995-04-24 00:00:00 1995-04
## 8 cs18         449552 0 0 1995-04-24 00:00:00 1995-04
## 9 cs18         456142 0 0 1995-04-24 00:00:00 1995-04
## 10 cs18        458942 0 0 1995-04-24 00:00:00 1995-04
## # ... with 147,025 more rows
```

```
# pull out one column, equivalent to using $:
```

```
downloads %>% pull(machineName) %>% head()
```

```
## [1] "cs18" "cs18" "cs18" "cs18" "cs18" "cs18"
```

Below is an example of how to use the family of `_join` function included in tidyverse. They are useful for combining two (or more) datasets, even if the sets only contain partial/subset of information.

```
# Join tibbles with subsets of information together:
```

```
# machineName and power rank
```

```
downloads5 <- tibble(machineName=c("cs18","piglet","tweetie","kermit", "pluto"),
                     powerRank=c(2,4,1,3,5))
```

```
# machineName and location of machine
```

```
downloads6 <- tibble(machineName=c("cs18","tweetie","kermit","skeeter"),
```

```
location=c("China", "USA", "Germany", "Japan"))
```

```
# all machineNames from tibble on the left are kept  
left_join(dowloads5, dowloads6)
```

```
## # A tibble: 5 x 3  
##   machineName powerRank location  
##   <chr>         <dbl> <chr>  
## 1 cs18           2 China  
## 2 piglet         4 <NA>  
## 3 tweetie        1 USA  
## 4 kermit         3 Germany  
## 5 pluto          5 <NA>
```

```
# all machineNames from tibble on the right are kept  
right_join(dowloads5, dowloads6)
```

```
## # A tibble: 4 x 3  
##   machineName powerRank location  
##   <chr>         <dbl> <chr>  
## 1 cs18           2 China  
## 2 tweetie        1 USA  
## 3 kermit         3 Germany  
## 4 skeeter        NA Japan
```

```
# only machineNames in both left and right tibble are kept  
inner_join(dowloads5, dowloads6)
```

```
## # A tibble: 3 x 3  
##   machineName powerRank location  
##   <chr>         <dbl> <chr>  
## 1 cs18           2 China  
## 2 tweetie        1 USA  
## 3 kermit         3 Germany
```

```
# all machineNames, from both tibbles are kept  
full_join(dowloads5, dowloads6)
```

```
## # A tibble: 6 x 3  
##   machineName powerRank location  
##   <chr>         <dbl> <chr>  
## 1 cs18           2 China  
## 2 piglet         4 <NA>  
## 3 tweetie        1 USA  
## 4 kermit         3 Germany  
## 5 pluto          5 <NA>  
## 6 skeeter        NA Japan
```