

II. Working with data in R (presentation)

Data Science Lab, University of Copenhagen

14 March, 2022

Tidyverse package

The tidyverse is a collection of R packages which, among other things, facilitate data handling and data transformation in R. See <https://www.tidyverse.org/> for details.

We must install and load the R package **tidyverse** before we have access to the functions.

- Install package: One option is to go via the *Tools* menu: *Tools* → *Install packages* → write **tidyverse** in the field called *Packages*. This only has to be done once. Otherwise use the **install.packages** function as shown here:

```
install.packages("tidyverse", repos = "https://mirrors.dotsrc.org/cran/")
```

- Load package: Use the **library** command below (preferred), or go to the *Packages* menu in the bottom right window, find **tidyverse** in the list, and click it. This has to be done in every R-session where you use the package.

```
library(tidyverse)
```

People with SCIENCE PC's (Windows) sometimes have problems with the installation step because R tries to install files to a place, where the user doesn't have permissions to save and edit files. You can try this instead:

- When you start RStudio, right-click the icon and choose *Run as administrator*. Perhaps you can now install packages by clicking *Tools* and *Include Packages* as above.
- If not, then the problem may be that RStudio is trying to install to your science drive (H: or \\a00143.science.domain). If so, try the command **.libPaths()**. If it shows two folders - one at the science drive and one locally one on your computer (C:) - then try the command **install.packages("tidyverse", lib=.libPaths()[2])**.

About the working directory

When working on a project, it is important to know *where you are*. The working directory is the path on your computer that R will *try* to access files from.

There are several helpful commands that help you navigate.

```
# show current working directory (cwd)
getwd()
```

```
# absolute path
setwd("~/Desktop/FromExceltoR/")
```

```
# relative path
```

```

setwd('./Presentations')

# go one step back in the directory
setwd('..')

# show folders in cwd
list.dirs(path = ".", recursive = FALSE)

# set working directory absolute path
setwd("~/Desktop/FromExceltoR/Presentations")

```

Import data

Data from Excel files can be imported via the *Import Dataset* facility. You may get the message that the package **readxl** should be installed. If so, then install it as explained for **tidyverse** above.

- Find *Import Data* in the upper right window in RStudio, and choose *From Excel* in the dropdown menu.
- A new window opens. Browse for the relevant Excel file; then a preview of the dataset is shown. Check that it looks OK, and click *Import*.
- Three things happened: Three lines of code was generated (and executed) in the Console, a new dataset now appears in the Environment window, and the dataset is shown in the top left window. Check again that it looks OK.
- Copy the first two lines of code into your R script (or into an R chunk in your Markdown document), but delete line starting with **View** and write instead the name of the dataset, here **downloads**. Then the first 10 lines of the data set are printed.

```

library(readxl)
downloads <- read_excel("downloads.xlsx")
downloads

```

```

## # A tibble: 147,035 x 6
##   machineName userID size time date month
##   <chr>      <dbl> <dbl> <dbl> <dtm> <chr>
## 1 cs18      146579 2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18      995988 7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18      317649 6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18      748501 13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18      955815 356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18      444174 0 0 1995-04-24 00:00:00 1995-04
## 7 cs18      446911 0 0 1995-04-24 00:00:00 1995-04
## 8 cs18      449552 0 0 1995-04-24 00:00:00 1995-04
## 9 cs18      456142 0 0 1995-04-24 00:00:00 1995-04
## 10 cs18     458942 0 0 1995-04-24 00:00:00 1995-04
## # ... with 147,025 more rows

```

R has stored the data in a so-called *tibble*, a type of data frame. Rows are referred to as *observations* or *data lines*, columns as *variables*. The data rows appear in the order as in the Excel file.

A slight digression: If data are saved in a csv file (comma separated values), possibly generated via an Excel sheet, then data can be read with the `read_csv` function. For example, if the data file is called `mydata.csv` and values are separated with commas, then the command

```
mydata <- read.csv("mydata.csv", sep=",")
```

creates a data frame in R with the data. The data frame is *not* a tibble and some of the commands below would not work for such a data frame.

About the data

The dataset is from Boston University and is about www data transfers from November 1994 to May 1995, see <http://ita.ee.lbl.gov/html/contrib/BU-Web-Client.html>.

- It has 147,035 data lines and 6 variables
 - *size* is the download size in bytes, and *time* is the download time in seconds.
-

Extracting variables, simple summary statistics

Variables can be extracted with the `$`-syntax, and we can use squared brackets to show only the first 40, say, values.

```
time_vector <- downloads$time  
time_vector[1:40]
```

```
## [1] 0.493030 0.325608 0.313704 0.582537 0.259252 0.000000 0.000000 0.000000  
## [9] 0.000000 0.000000 0.000000 0.335502 0.284853 0.000000 0.000000 0.000000  
## [17] 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000  
## [25] 0.285665 0.397111 3.410561 0.267474 0.842364 0.903005 2.784645 2.806157  
## [33] 0.990092 0.477629 0.000000 0.000000 0.000000 0.000000 0.988944 0.000000
```

Summary statistics like mean, standard deviation, median are easily computed for a vector.

Examples of R functions for computing summary statistics: `length`, `mean`, `median`, `sd`, `var`, `sum`, `quantile`, `min`, `max`, `IQR`.

```
length(time_vector)
```

```
## [1] 147035
```

```
mean(time_vector)
```

```
## [1] 0.9539674
```

```
sd(time_vector)
```

```
## [1] 14.22557
```

```
median(time_vector)
```

```
## [1] 0
```

```
min(time_vector)
```

```
## [1] 0
```

Notice that more than half the observations have time equal to zero (median is zero).

Filtering data (selecting rows): filter

The `filter` function is used to make sub-datasets where only certain datalines (rows) are maintained. It's described with *logical expressions* which datalines should be kept in the dataset.

Say that we only want observations with download time larger than 1000 seconds; there happens to be eight such observations:

```
filter(downloads, time > 1000)
```

```
## # A tibble: 8 x 6
##   machineName userID      size time date      month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>    <chr>
## 1 cs18      502807  4055821 1275. 1994-12-02 00:00:00 1994-12
## 2 cs18      16653   2573336 1335. 1994-11-22 00:00:00 1994-11
## 3 cs18      957883   2743516 1151. 1994-11-22 00:00:00 1994-11
## 4 cs18       47910   4720220 1749. 1994-11-22 00:00:00 1994-11
## 5 tweetie   223655    245003 1214. 1995-04-13 00:00:00 1995-04
## 6 kermit    576790  14518894 1380. 1995-04-20 00:00:00 1995-04
## 7 kermit    139654   1079731 1129. 1995-02-23 00:00:00 1995-02
## 8 pluto     337530   8674562 1878. 1995-03-13 00:00:00 1995-03
```

```
downloads %>%
  filter(time > 1000)
```

```
## # A tibble: 8 x 6
##   machineName userID      size time date      month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>    <chr>
## 1 cs18      502807  4055821 1275. 1994-12-02 00:00:00 1994-12
## 2 cs18      16653   2573336 1335. 1994-11-22 00:00:00 1994-11
## 3 cs18      957883   2743516 1151. 1994-11-22 00:00:00 1994-11
## 4 cs18       47910   4720220 1749. 1994-11-22 00:00:00 1994-11
## 5 tweetie   223655    245003 1214. 1995-04-13 00:00:00 1995-04
## 6 kermit    576790  14518894 1380. 1995-04-20 00:00:00 1995-04
## 7 kermit    139654   1079731 1129. 1995-02-23 00:00:00 1995-02
## 8 pluto     337530   8674562 1878. 1995-03-13 00:00:00 1995-03
```

Or say that only want observations with strictly positive download size:

```
downloads2 <- filter(downloads, size > 0)
downloads2
```

```
## # A tibble: 36,708 x 6
##   machineName userID  size time date      month
##   <chr>      <dbl> <dbl> <dbl> <dtm>    <chr>
## 1 cs18      146579   2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18      995988   7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18      317649   6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18      748501  13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18      955815    356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18      596819  15063 0.336 1995-04-24 00:00:00 1995-04
## 7 cs18      169424   2548 0.285 1995-04-24 00:00:00 1995-04
## 8 cs18      386686   1932 0.286 1995-04-24 00:00:00 1995-04
## 9 cs18      783767   7294 0.397 1995-04-24 00:00:00 1995-04
## 10 cs18     788633   4470 3.41 1995-04-24 00:00:00 1995-04
## # ... with 36,698 more rows
```

Notice that this result is assigned to **downloads2**. It has 36,708 data lines. The original data called

`downloads` still exists with 147,035 data lines.

Filtering requires *logical predicates*. These are expressions in terms of columns, which evaluate to either TRUE or FALSE for each row. Logical expressions can be combined with logical operations.

- Comparisons: `==`, `!=`, `<`, `>`, `<=`, `>=`, `%in%`, `is.na`
- Logical operations: `!` (not), `|` (or), `&` (and). A comma can be used instead of `&`

Here comes two sub-datasets:

```
# Rows from kermit, and with size greater than 200000 bytes are kept.
filter(downloads2, machineName == "kermit", size > 200000)
```

```
## # A tibble: 98 x 6
##   machineName userID      size    time date      month
##   <chr>      <dbl>    <dbl>    <dbl> <dtm>      <chr>
## 1 kermit    157161  498325  0.629 1995-04-13 00:00:00 1995-04
## 2 kermit    734988  271058  17.3   1995-04-22 00:00:00 1995-04
## 3 kermit    388066  435923  29.2   1995-04-22 00:00:00 1995-04
## 4 kermit     34030  642771  4.80   1995-04-12 00:00:00 1995-04
## 5 kermit    327021  724757  4.98   1995-04-12 00:00:00 1995-04
## 6 kermit     38016  561762  9.75   1995-04-05 00:00:00 1995-04
## 7 kermit    277395  404209  11.3   1995-04-05 00:00:00 1995-04
## 8 kermit    576790 14518894 1380.   1995-04-20 00:00:00 1995-04
## 9 kermit     17623  489473  21.2   1995-04-20 00:00:00 1995-04
## 10 kermit   198041  355963  15.3   1995-04-20 00:00:00 1995-04
## # ... with 88 more rows
```

```
# Rows NOT from kermit, and with size greater than 200000 bytes are kept.
filter(downloads2, machineName != "kermit" & size > 200000)
```

```
## # A tibble: 220 x 6
##   machineName userID      size    time date      month
##   <chr>      <dbl>    <dbl>    <dbl> <dtm>      <chr>
## 1 cs18      204764 2691689  0.834 1995-04-26 00:00:00 1995-04
## 2 cs18      397405  215045  1.10   1994-12-15 00:00:00 1994-12
## 3 cs18      809091  226586  3.92   1994-12-15 00:00:00 1994-12
## 4 cs18      779032 1080472 156.    1994-12-11 00:00:00 1994-12
## 5 cs18      688294  748705  93.1   1994-12-11 00:00:00 1994-12
## 6 cs18      447740 6360764 863.    1994-12-11 00:00:00 1994-12
## 7 cs18      708452  204918  7.07   1994-12-18 00:00:00 1994-12
## 8 cs18      598668  204918  12.7   1994-12-18 00:00:00 1994-12
## 9 cs18      288167  204918  4.98   1994-12-18 00:00:00 1994-12
## 10 cs18     974956  203714  6.13   1994-12-16 00:00:00 1994-12
## # ... with 210 more rows
```

A helpful function to know which machine names are valid can be:

```
# get unique machineName values in downloads2
distinct(downloads2, machineName)
```

```
## # A tibble: 5 x 1
##   machineName
##   <chr>
## 1 cs18
## 2 piglet
## 3 kermit
## 4 tweetie
```

```
## 5 pluto
```

And if you are looking for multiple values for a given variable:

```
downloads2 %>% filter(machineName %in% c("kermit","pluto"), size > 2000000)
```

```
## # A tibble: 8 x 6
##   machineName userID      size    time date      month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>    <chr>
## 1 kermit      576790 14518894 1380. 1995-04-20 00:00:00 1995-04
## 2 kermit      756949  4418124  439. 1995-04-20 00:00:00 1995-04
## 3 kermit      287308  6935603   88.2 1995-04-24 00:00:00 1995-04
## 4 kermit      928227  9523767  171. 1995-02-08 00:00:00 1995-02
## 5 kermit      128147  2743816  216. 1995-02-23 00:00:00 1995-02
## 6 pluto       867173  4670973  230. 1995-03-14 00:00:00 1995-03
## 7 kermit      456524  2836135  127. 1995-03-31 00:00:00 1995-03
## 8 pluto       337530  8674562 1878. 1995-03-13 00:00:00 1995-03
```

Selecting variables: `select`

Sometimes, datasets has many variables of which only some are relevant for the analysis. Variables can be selected or skipped with the `select` function.

```
# Without the date variable
```

```
select(downloads2, -date)
```

```
## # A tibble: 36,708 x 5
##   machineName userID  size  time month
##   <chr>      <dbl> <dbl> <dbl> <chr>
## 1 cs18      146579  2464 0.493 1995-04
## 2 cs18      995988  7745 0.326 1995-04
## 3 cs18      317649  6727 0.314 1995-04
## 4 cs18      748501 13049 0.583 1995-04
## 5 cs18      955815   356 0.259 1995-04
## 6 cs18      596819 15063 0.336 1995-04
## 7 cs18      169424  2548 0.285 1995-04
## 8 cs18      386686  1932 0.286 1995-04
## 9 cs18      783767  7294 0.397 1995-04
## 10 cs18     788633  4470 3.41 1995-04
## # ... with 36,698 more rows
```

```
# Only include the three mentioned variable names
```

```
downloads3 <- select(downloads2, machineName, size, time)
downloads3
```

```
## # A tibble: 36,708 x 3
##   machineName  size  time
##   <chr>      <dbl> <dbl>
## 1 cs18      2464 0.493
## 2 cs18      7745 0.326
## 3 cs18      6727 0.314
## 4 cs18     13049 0.583
## 5 cs18       356 0.259
## 6 cs18     15063 0.336
## 7 cs18      2548 0.285
## 8 cs18      1932 0.286
```

```
## 9 cs18          7294 0.397
## 10 cs18         4470 3.41
## # ... with 36,698 more rows
```

Notice that we have made a new dataframe, **downloads3** with only three variables.

Transformations of data

Transformations of existing variables in the data set can be computed and included in the data set with the `mutate` function.

We first compute two new variables, download speed (**speed**) and the logarithm of the download size (**logSize**):

```
downloads3 <- mutate(downloads3, speed = size / time, logSize = log10(size))
downloads3
```

```
## # A tibble: 36,708 x 5
##   machineName size time speed logSize
##   <chr>      <dbl> <dbl> <dbl> <dbl>
## 1 cs18        2464 0.493 4998.   3.39
## 2 cs18        7745 0.326 23786.  3.89
## 3 cs18        6727 0.314 21444.  3.83
## 4 cs18       13049 0.583 22400.  4.12
## 5 cs18         356 0.259 1373.   2.55
## 6 cs18       15063 0.336 44897.  4.18
## 7 cs18        2548 0.285 8945.   3.41
## 8 cs18        1932 0.286 6763.   3.29
## 9 cs18        7294 0.397 18368.  3.86
## 10 cs18       4470 3.41 1311.   3.65
## # ... with 36,698 more rows
```

We then make a new categorical variable, **slow**, which is “Yes” if speed < 150 and “No” otherwise

```
downloads3 <- mutate(downloads3, slow = ifelse(speed < 150, "Yes", "No"))
downloads3
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18        2464 0.493 4998.   3.39 No
## 2 cs18        7745 0.326 23786.  3.89 No
## 3 cs18        6727 0.314 21444.  3.83 No
## 4 cs18       13049 0.583 22400.  4.12 No
## 5 cs18         356 0.259 1373.   2.55 No
## 6 cs18       15063 0.336 44897.  4.18 No
## 7 cs18        2548 0.285 8945.   3.41 No
## 8 cs18        1932 0.286 6763.   3.29 No
## 9 cs18        7294 0.397 18368.  3.86 No
## 10 cs18       4470 3.41 1311.   3.65 No
## # ... with 36,698 more rows
```

Counting, tabulation of categorical variables: `count`

The `count` function is useful for counting data datalines, possibly according to certain criteria or for the different levels of categorical values.

```
# Total number of observations in the current dataset
```

```
count(downloads3)
```

```
## # A tibble: 1 x 1
```

```
##       n
```

```
##   <int>
```

```
## 1 36708
```

```
# Number of observations from each machine
```

```
count(downloads3, machineName)
```

```
## # A tibble: 5 x 2
```

```
##   machineName      n
```

```
##   <chr>         <int>
```

```
## 1 cs18          3814
```

```
## 2 kermit        9094
```

```
## 3 piglet       11200
```

```
## 4 pluto        5253
```

```
## 5 tweetie      7347
```

```
# Number of observations which have/have not size larger than 5000
```

```
count(downloads3, size>5000)
```

```
## # A tibble: 2 x 2
```

```
##   `size > 5000`      n
```

```
##   <lgl>          <int>
```

```
## 1 FALSE         25865
```

```
## 2 TRUE          10843
```

```
# Number of observations for each combination of machine name and the *slow* variable.
```

```
count(downloads3, machineName, slow)
```

```
## # A tibble: 10 x 3
```

```
##   machineName slow      n
```

```
##   <chr>         <chr> <int>
```

```
## 1 cs18         No     3662
```

```
## 2 cs18         Yes     152
```

```
## 3 kermit       No     8717
```

```
## 4 kermit       Yes     377
```

```
## 5 piglet       No    10734
```

```
## 6 piglet       Yes     466
```

```
## 7 pluto        No    4963
```

```
## 8 pluto        Yes     290
```

```
## 9 tweetie      No    6983
```

```
## 10 tweetie     Yes     364
```

Sorting data: `arrange`

The `arrange` function can be used to sort the data according to one or more columns.

Let's sort the data according to download size (ascending order). The first lines of the sorted data set is printed on-screen, but the dataset `downloads3` has *not* been changed.


```
arrange(downloads3, size)
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18         3  3.73 0.804  0.477 Yes
## 2 piglet       3  1.53 1.96   0.477 Yes
## 3 piglet       3  1.53 1.96   0.477 Yes
## 4 tweetie      3  1.11 2.71   0.477 Yes
## 5 kermit       3  1.12 2.69   0.477 Yes
## 6 pluto        3  8.60 0.349  0.477 Yes
## 7 pluto        3  9.87 0.304  0.477 Yes
## 8 pluto        3  3.78 0.793  0.477 Yes
## 9 pluto        3  4.68 0.641  0.477 Yes
##10 pluto        3  4.93 0.608  0.477 Yes
## # ... with 36,698 more rows
```

Two different examples:

```
# According to download size in descending order
arrange(downloads3, desc(size))
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 kermit    14518894 1380. 10522.  7.16 No
## 2 piglet    14158123 123. 115169.  7.15 No
## 3 kermit     9523767 171. 55562.  6.98 No
## 4 piglet     9384067 80.0 117309.  6.97 No
## 5 pluto      8674562 1878. 4619.  6.94 No
## 6 kermit     6935603 88.2 78655.  6.84 No
## 7 cs18       6360764 863. 7374.  6.80 No
## 8 piglet     5143062 597. 8611.  6.71 No
## 9 piglet     4812334 215. 22345.  6.68 No
##10 cs18       4720220 1749. 2700.  6.67 No
## # ... with 36,698 more rows
```

```
# After machine name and then according to download size in descending order
arrange(downloads3, machineName, desc(size))
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18       6360764 863. 7374.  6.80 No
## 2 cs18       4720220 1749. 2700.  6.67 No
## 3 cs18       4055821 1275. 3180.  6.61 No
## 4 cs18       3047343 20.9 146038.  6.48 No
## 5 cs18       2952381 318. 9289.  6.47 No
## 6 cs18       2743516 1151. 2383.  6.44 No
## 7 cs18       2691689 0.834 3228695.  6.43 No
## 8 cs18       2613025 18.5 140959.  6.42 No
## 9 cs18       2573336 1335. 1928.  6.41 No
##10 cs18       1931453 186. 10388.  6.29 No
## # ... with 36,698 more rows
```

Grouping: `group_by`

We can group the dataset by one or more categorical variables with `group_by`. The dataset is not changed as such, but - as we will see - grouping can be useful for computation of summary statistics and graphics.

Here we group after machine name (first) *and* the slow variable (second). The only way we can see it at this point is in the second line in the output (`# Groups:`):

```
# Group according to machine
group_by(downloads3, machineName)

## # A tibble: 36,708 x 6
## # Groups:   machineName [5]
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18        2464 0.493 4998.    3.39 No
## 2 cs18        7745 0.326 23786.   3.89 No
## 3 cs18        6727 0.314 21444.   3.83 No
## 4 cs18       13049 0.583 22400.   4.12 No
## 5 cs18         356 0.259 1373.    2.55 No
## 6 cs18       15063 0.336 44897.   4.18 No
## 7 cs18        2548 0.285 8945.    3.41 No
## 8 cs18        1932 0.286 6763.    3.29 No
## 9 cs18        7294 0.397 18368.   3.86 No
## 10 cs18       4470 3.41 1311.    3.65 No
## # ... with 36,698 more rows
```

```
# Group according to machine and slow
group_by(downloads3, machineName, slow)

## # A tibble: 36,708 x 6
## # Groups:   machineName, slow [10]
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18        2464 0.493 4998.    3.39 No
## 2 cs18        7745 0.326 23786.   3.89 No
## 3 cs18        6727 0.314 21444.   3.83 No
## 4 cs18       13049 0.583 22400.   4.12 No
## 5 cs18         356 0.259 1373.    2.55 No
## 6 cs18       15063 0.336 44897.   4.18 No
## 7 cs18        2548 0.285 8945.    3.41 No
## 8 cs18        1932 0.286 6763.    3.29 No
## 9 cs18        7294 0.397 18368.   3.86 No
## 10 cs18       4470 3.41 1311.    3.65 No
## # ... with 36,698 more rows
```

Summary statistics, revisited: `summarize`

Recall how we could compute summary statistics for a single variable in a dataset, e.g.

```
mean(downloads3$size)
```

```
## [1] 16638.36
```

```
max(downloads3$size)
```

```
## [1] 14518894
```

With `summarize` we can compute summary statistics for a variable for each level of a grouping variable or for each combination of several grouping variables.

First, a bunch of summaries for the size variable for each machine name, where we give explicit names for the new variables:

```
downloads.grp1 <- group_by(downloads3, machineName)
summarize(downloads.grp1,
  avg = mean(size),
  med = median(size),
  stdev = sd(size),
  total = sum(size),
  n = n())
```

```
## # A tibble: 5 x 6
##   machineName   avg   med   stdev   total     n
##   <chr>       <dbl> <dbl>   <dbl>   <dbl> <int>
## 1 cs18       26375. 1990. 208915. 100593281 3814
## 2 kermit     19247. 2466 213985. 175032552 9094
## 3 piglet     14121. 2146 188340. 158149841 11200
## 4 pluto      13822. 2069 144425. 72605544 5253
## 5 tweetie    14207. 2197 94318. 104379794 7347
```

Second, the same thing but for each combination of machine name and the slow variable:

```
downloads.grp2 <- group_by(downloads3, machineName, slow)
summarize(downloads.grp2,
  avg = mean(size),
  med = median(size),
  stdev = sd(size),
  total = sum(size),
  n = n())
```

```
## # A tibble: 10 x 7
## # Groups:   machineName [5]
##   machineName slow   avg   med   stdev   total     n
##   <chr>       <chr> <dbl> <dbl>   <dbl>   <dbl> <int>
## 1 cs18       No    27445. 2092. 213140. 100503042 3662
## 2 cs18       Yes     594. 368.   614.    90239    152
## 3 kermit     No    20030. 2598 218529. 174602282 8717
## 4 kermit     Yes    1141. 541   3049.   430270    377
## 5 piglet     No    14687. 2264 192365. 157650747 10734
## 6 piglet     Yes    1071. 416.   1934.   499094    466
## 7 pluto      No    14564. 2164 148551. 72280790 4963
## 8 pluto      Yes    1120. 413   2108.   324754    290
## 9 tweetie    No    14894. 2373 96694. 104001733 6983
## 10 tweetie   Yes    1039. 471   2603.   378061    364
```

Third, mean and standard deviation for several variables:

```
summarize_at(downloads.grp2, c("time", "size"), list(ave=mean, stdev=sd))
```

```
## # A tibble: 10 x 6
## # Groups:   machineName [5]
##   machineName slow  time_ave size_ave time_stdev size_stdev
##   <chr>       <chr>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 cs18       No      5.17   27445.    57.1   213140.
## 2 cs18       Yes     9.63    594.    17.8    614.
```

##	3	kermit	No	3.41	20030.	25.3	218529.
##	4	kermit	Yes	20.7	1141.	47.8	3049.
##	5	piglet	No	2.33	14687.	13.8	192365.
##	6	piglet	Yes	19.4	1071.	40.2	1934.
##	7	pluto	No	3.40	14564.	30.4	148551.
##	8	pluto	Yes	21.7	1120.	46.3	2108.
##	9	tweetie	No	2.68	14894.	17.3	96694.
##	10	tweetie	Yes	17.8	1039.	34.5	2603.

The datasets with summaries can be saved as datasets themselves, for example to be used as the basis for certain graphs.

The pipe operator: %>%

Two or more function calls can be evaluated sequentially using the so-called pipe operator, %>%. Nesting of function calls becomes more readable, and intermediate assignments are avoided.

Let's try it to do a bunch of things in one go, starting with the original dataset:

```
downloads %>%
  filter(size>0) %>% # Subset of data
  group_by(machineName) %>% # Grouping
  summarize(avg = mean(size)) %>% # Compute mean
  arrange(avg) # Sort after mean
```

```
## # A tibble: 5 x 2
##   machineName    avg
##   <chr>         <dbl>
## 1 pluto         13822.
## 2 piglet        14121.
## 3 tweetie       14207.
## 4 kermit        19247.
## 5 cs18          26375.
```