

## II. Working with data in R (presentation)

Data Science Lab, University of Copenhagen

18 October, 2021

### Tidyverse package

The tidyverse is a collection of R packages which, among other things, facilitate data handling and data transformation in R. See <https://www.tidyverse.org/> for details.

We must install and load the R package **tidyverse** before we have access to the functions.

- Install package: One option is to go via the *Tools* menu: *Tools* → *Install packages* → write **tidyverse** in the field called *Packages*. This only has to be done once. Otherwise use the **install.packages** function as shown here:

```
install.packages("tidyverse", repos = "https://mirrors.dotsrc.org/cran/")
```

```
##
```

```
## The downloaded binary packages are in
```

```
## /var/folders/ch/mtlydjm96n7cqq2d3b3xjdw0000gp/T/RtmpV92mVM/downloaded_packages
```

- Load package: Use the **library** command below (preferred), or go to the *Packages* menu in the bottom right window, find **tidyverse** in the list, and click it. This has to be done in every R-session where you use the package.

```
library(tidyverse)
```

People with SCIENCE PC's (Windows) sometimes have problems with the installation step because R tries to install files to a place, where the user doesn't have permissions to save and edit files. You can try this instead:

- When you start RStudio, right-click the icon and choose *Run as administrator*. Perhaps you can now install packages by clicking *Tools* and *Include Packages* as above.
- If not, then the problem may be that RStudio is trying to install to your science drive (H: or \\a00143.science.domain). If so, try the command **.libPaths()**. If it shows two folders - one at the science drive and one locally one on your computer (C:) - then try the command **install.packages("tidyverse", lib=.libPaths()[2])**.

---

### About the working directory

When working on a project, it is important to know *where you are*. The working directory is the path on your computer that R will *try* to access files from.

There are several helpful commands that help you navigate.

```
# show current working directory (cwd)
getwd()
```

```
# show folders in cwd
```

```
list.dirs(path = ".", recursive = FALSE)

# relative path, go one folder deeper into Presentations
setwd('./Presentations')

# go one step back in the directory
setwd('..')

# absolute path
setwd("~/Documents/work/Projects/FromExceltoR_2021/Presentations")
```

## Import data

Data from Excel files can be imported via the *Import Dataset* facility. You may get the message that the package **readxl** should be installed. If so, then install it as explained for **tidyverse** above.

- Find *Import Data* in the upper right window in RStudio, and choose *From Excel* in the dropdown menu.
- A new window opens. Browse for the relevant Excel file; then a preview of the dataset is shown. Check that it looks OK, and click *Import*.
- Three things happened: Three lines of code was generated (and executed) in the Console, a new dataset now appears in the Environment window, and the dataset is shown in the top left window. Check again that it looks OK.
- Copy the first two lines of code into your R script (or into an R chunk in your Markdown document), but delete line starting with **View** and write instead the name of the dataset, here **downloads**. Then the first 10 lines of the data set are printed.

```
library(readxl)
downloads <- read_excel("downloads.xlsx")
downloads
```

```
## # A tibble: 147,035 x 6
##   machineName userID size time date month
##   <chr>      <dbl> <dbl> <dbl> <dtm> <chr>
## 1 cs18      146579 2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18      995988 7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18      317649 6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18      748501 13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18      955815 356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18      444174 0 0 1995-04-24 00:00:00 1995-04
## 7 cs18      446911 0 0 1995-04-24 00:00:00 1995-04
## 8 cs18      449552 0 0 1995-04-24 00:00:00 1995-04
## 9 cs18      456142 0 0 1995-04-24 00:00:00 1995-04
## 10 cs18     458942 0 0 1995-04-24 00:00:00 1995-04
## # ... with 147,025 more rows
```

R has stored the data in a so-called *tibble*, a type of data frame. Rows are referred to as *observations* or *data lines*, columns as *variables*. The data rows appear in the order as in the Excel file.

A slight digression: If data are saved in a csv file (comma separated values), possibly generated via an Excel sheet, then data can be read with the `read_csv` function. For example, if the data file is called `mydata.csv` and values are separated with commas, then the command

```
mydata <- read.csv("mydata.csv", sep=",")
```

creates a data frame in R with the data. The data frame is *not* a tibble and some of the commands below would not work for such a data frame.

---

## About the data

The dataset is from Boston University and is about www data transfers from November 1994 to May 1995, see <http://ita.ee.lbl.gov/html/contrib/BU-Web-Client.html>.

- It has 147,035 data lines and 6 variables
  - *size* is the download size in bytes, and *time* is the download time in seconds.
- 

## Extracting variables, simple summary statistics

Variables can be extracted with the `$`-syntax, and we can use squared brackets to show only the first 40, say, values.

```
time_vector <- downloads$time  
time_vector[1:40]
```

```
## [1] 0.493030 0.325608 0.313704 0.582537 0.259252 0.000000 0.000000 0.000000  
## [9] 0.000000 0.000000 0.000000 0.335502 0.284853 0.000000 0.000000 0.000000  
## [17] 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000  
## [25] 0.285665 0.397111 3.410561 0.267474 0.842364 0.903005 2.784645 2.806157  
## [33] 0.990092 0.477629 0.000000 0.000000 0.000000 0.000000 0.988944 0.000000
```

Summary statistics like mean, standard deviation, median are easily computed for a vector.

Examples of R functions for computing summary statistics: `length`, `mean`, `median`, `sd`, `var`, `sum`, `quantile`, `min`, `max`, `IQR`.

```
length(time_vector)
```

```
## [1] 147035
```

```
mean(time_vector)
```

```
## [1] 0.9539674
```

```
sd(time_vector)
```

```
## [1] 14.22557
```

```
median(time_vector)
```

```
## [1] 0
```

```
min(time_vector)
```

```
## [1] 0
```

Notice that more than half the observations have time equal to zero (median is zero).

---

## Filtering data (selecting rows): `filter`

The `filter` function is used to make sub-datasets where only certain datalines (rows) are maintained. It's described with *logical expressions* which datalines should be kept in the dataset.

Say that we only want observations with download time larger than 1000 seconds; there happens to be eight such observations:

```
filter(downloads, time > 1000)
```

```
## # A tibble: 8 x 6
##   machineName userID      size time date      month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>      <chr>
## 1 cs18      502807  4055821 1275. 1994-12-02 00:00:00 1994-12
## 2 cs18      16653   2573336 1335. 1994-11-22 00:00:00 1994-11
## 3 cs18      957883   2743516 1151. 1994-11-22 00:00:00 1994-11
## 4 cs18       47910   4720220 1749. 1994-11-22 00:00:00 1994-11
## 5 tweetie   223655    245003 1214. 1995-04-13 00:00:00 1995-04
## 6 kermit    576790  14518894 1380. 1995-04-20 00:00:00 1995-04
## 7 kermit    139654   1079731 1129. 1995-02-23 00:00:00 1995-02
## 8 pluto     337530   8674562 1878. 1995-03-13 00:00:00 1995-03
```

Or say that only want observations with strictly positive download size:

```
downloads2 <- filter(downloads, size > 0)
downloads2
```

```
## # A tibble: 36,708 x 6
##   machineName userID      size time date      month
##   <chr>      <dbl>    <dbl> <dbl> <dtm>      <chr>
## 1 cs18      146579    2464 0.493 1995-04-24 00:00:00 1995-04
## 2 cs18      995988    7745 0.326 1995-04-24 00:00:00 1995-04
## 3 cs18      317649    6727 0.314 1995-04-24 00:00:00 1995-04
## 4 cs18      748501   13049 0.583 1995-04-24 00:00:00 1995-04
## 5 cs18      955815     356 0.259 1995-04-24 00:00:00 1995-04
## 6 cs18      596819   15063 0.336 1995-04-24 00:00:00 1995-04
## 7 cs18      169424    2548 0.285 1995-04-24 00:00:00 1995-04
## 8 cs18      386686    1932 0.286 1995-04-24 00:00:00 1995-04
## 9 cs18      783767    7294 0.397 1995-04-24 00:00:00 1995-04
## 10 cs18     788633    4470 3.41 1995-04-24 00:00:00 1995-04
## # ... with 36,698 more rows
```

Notice that this result is assigned to **downloads2**. It has 36,708 data lines. The original data called **downloads** still exists with 147,035 data lines.

Filtering requires *logical predicates*. These are expressions in terms of columns, which evaluate to either TRUE or FALSE for each row. Logical expressions can be combined with logical operations.

- Comparisons: `==`, `!=`, `<`, `>`, `<=`, `>=`, `%in%`, `is.na`
- Logical operations: `!` (not), `|` (or), `&` (and). A comma can be used instead of `&`

Here comes two sub-datasets:

```
# Rows from kermit, and with size greater than 200000 bytes are kept.
filter(downloads2, machineName == "kermit", size > 200000)
```

```
## # A tibble: 98 x 6
##   machineName userID      size      time date      month
##   <chr>      <dbl>    <dbl>    <dbl> <dtm>      <chr>
```

```
## 1 kermit      157161  498325    0.629 1995-04-13 00:00:00 1995-04
## 2 kermit      734988  271058    17.3  1995-04-22 00:00:00 1995-04
## 3 kermit      388066  435923    29.2  1995-04-22 00:00:00 1995-04
## 4 kermit       34030  642771    4.80  1995-04-12 00:00:00 1995-04
## 5 kermit      327021  724757    4.98  1995-04-12 00:00:00 1995-04
## 6 kermit       38016  561762    9.75  1995-04-05 00:00:00 1995-04
## 7 kermit      277395  404209   11.3  1995-04-05 00:00:00 1995-04
## 8 kermit      576790 14518894 1380.  1995-04-20 00:00:00 1995-04
## 9 kermit       17623  489473   21.2  1995-04-20 00:00:00 1995-04
## 10 kermit     198041  355963   15.3  1995-04-20 00:00:00 1995-04
## # ... with 88 more rows
```

```
# Rows NOT from kermit, and with size greater than 200000 bytes are kept.
filter(downloads2, machineName != "kermit" & size > 200000)
```

```
## # A tibble: 220 x 6
##   machineName userID      size    time date      month
##   <chr>      <dbl>    <dbl>    <dbl> <dtm>      <chr>
## 1 cs18      204764 2691689    0.834 1995-04-26 00:00:00 1995-04
## 2 cs18      397405 215045    1.10  1994-12-15 00:00:00 1994-12
## 3 cs18      809091 226586    3.92  1994-12-15 00:00:00 1994-12
## 4 cs18      779032 1080472 156.  1994-12-11 00:00:00 1994-12
## 5 cs18      688294 748705   93.1  1994-12-11 00:00:00 1994-12
## 6 cs18      447740 6360764 863.  1994-12-11 00:00:00 1994-12
## 7 cs18      708452 204918    7.07  1994-12-18 00:00:00 1994-12
## 8 cs18      598668 204918   12.7  1994-12-18 00:00:00 1994-12
## 9 cs18      288167 204918    4.98  1994-12-18 00:00:00 1994-12
## 10 cs18     974956 203714    6.13  1994-12-16 00:00:00 1994-12
## # ... with 210 more rows
```

A helpful function to know which machine names are valid can be:

```
# get unique machineName values in downloads2
distinct(downloads2, machineName)
```

```
## # A tibble: 5 x 1
##   machineName
##   <chr>
## 1 cs18
## 2 piglet
## 3 kermit
## 4 tweetie
## 5 pluto
```

And if you are looking for multiple values for a given variable:

```
downloads2 %>% filter(machineName %in% c("kermit", "pluto"), size > 2000000)
```

```
## # A tibble: 8 x 6
##   machineName userID      size    time date      month
##   <chr>      <dbl>    <dbl>    <dbl> <dtm>      <chr>
## 1 kermit      576790 14518894 1380.  1995-04-20 00:00:00 1995-04
## 2 kermit      756949 4418124  439.  1995-04-20 00:00:00 1995-04
## 3 kermit      287308 6935603   88.2  1995-04-24 00:00:00 1995-04
## 4 kermit      928227 9523767  171.  1995-02-08 00:00:00 1995-02
## 5 kermit      128147 2743816  216.  1995-02-23 00:00:00 1995-02
## 6 pluto       867173 4670973  230.  1995-03-14 00:00:00 1995-03
```

```
## 7 kermi      456524 2836135 127. 1995-03-31 00:00:00 1995-03
## 8 pluto      337530 8674562 1878. 1995-03-13 00:00:00 1995-03
```

---

## Selecting variables: `select`

Sometimes, datasets has many variables of which only some are relevant for the analysis. Variables can be selected or skipped with the `select` function.

```
# Without the date variable
```

```
select(downloads2, -date)
```

```
## # A tibble: 36,708 x 5
##   machineName userID size time month
##   <chr>      <dbl> <dbl> <dbl> <chr>
## 1 cs18      146579 2464 0.493 1995-04
## 2 cs18      995988 7745 0.326 1995-04
## 3 cs18      317649 6727 0.314 1995-04
## 4 cs18      748501 13049 0.583 1995-04
## 5 cs18      955815 356 0.259 1995-04
## 6 cs18      596819 15063 0.336 1995-04
## 7 cs18      169424 2548 0.285 1995-04
## 8 cs18      386686 1932 0.286 1995-04
## 9 cs18      783767 7294 0.397 1995-04
## 10 cs18     788633 4470 3.41 1995-04
## # ... with 36,698 more rows
```

```
# Only include the three mentioned variable names
```

```
downloads3 <- select(downloads2, machineName, size, time)
downloads3
```

```
## # A tibble: 36,708 x 3
##   machineName size time
##   <chr>      <dbl> <dbl>
## 1 cs18      2464 0.493
## 2 cs18      7745 0.326
## 3 cs18      6727 0.314
## 4 cs18     13049 0.583
## 5 cs18      356 0.259
## 6 cs18     15063 0.336
## 7 cs18      2548 0.285
## 8 cs18      1932 0.286
## 9 cs18      7294 0.397
## 10 cs18     4470 3.41
## # ... with 36,698 more rows
```

Notice that we have made a new dataframe, **downloads3** with only three variables.

---

## Transformations of data

Tranformations of existing variables in the data set can be computed and included in the data set with the `mutate` function.

We first compute two new variables, download speed (**speed**) and the logarithm of the download size (**logSize**):

```
downloads3 <- mutate(downloads3, speed = size / time, logSize = log10(size))
downloads3
```

```
## # A tibble: 36,708 x 5
##   machineName size time speed logSize
##   <chr>      <dbl> <dbl> <dbl> <dbl>
## 1 cs18      2464 0.493 4998. 3.39
## 2 cs18      7745 0.326 23786. 3.89
## 3 cs18      6727 0.314 21444. 3.83
## 4 cs18     13049 0.583 22400. 4.12
## 5 cs18       356 0.259 1373. 2.55
## 6 cs18     15063 0.336 44897. 4.18
## 7 cs18      2548 0.285 8945. 3.41
## 8 cs18      1932 0.286 6763. 3.29
## 9 cs18      7294 0.397 18368. 3.86
## 10 cs18     4470 3.41 1311. 3.65
## # ... with 36,698 more rows
```

We then make a new categorial variable, **slow**, which is “Yes” is speed < 150 and “No” otherwise

```
downloads3 <- mutate(downloads3, slow = ifelse(speed < 150, "Yes", "No"))
downloads3
```

```
## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18      2464 0.493 4998. 3.39 No
## 2 cs18      7745 0.326 23786. 3.89 No
## 3 cs18      6727 0.314 21444. 3.83 No
## 4 cs18     13049 0.583 22400. 4.12 No
## 5 cs18       356 0.259 1373. 2.55 No
## 6 cs18     15063 0.336 44897. 4.18 No
## 7 cs18      2548 0.285 8945. 3.41 No
## 8 cs18      1932 0.286 6763. 3.29 No
## 9 cs18      7294 0.397 18368. 3.86 No
## 10 cs18     4470 3.41 1311. 3.65 No
## # ... with 36,698 more rows
```

---

## Counting, tabulation of categorical variables: count

The `count` function is useful for counting data datalines, possibly according to certain criteria or for the different levels of categorical values.

```
# Total number of observations in the current dataset
count(downloads3)
```

```
## # A tibble: 1 x 1
##       n
##   <int>
## 1 36708
```

```
# Number of observations from each machine
count(downloads3, machineName)
```

```
## # A tibble: 5 x 2
##   machineName      n
```

```
##   <chr>      <int>
## 1 cs18      3814
## 2 kermit    9094
## 3 piglet    11200
## 4 pluto     5253
## 5 tweetie   7347

# Number of observations which have/have not size larger than 5000
count(downloads3, size>5000)

## # A tibble: 2 x 2
##   `size > 5000`      n
##   <lgl>          <int>
## 1 FALSE         25865
## 2 TRUE          10843

# Number of observations for each combination of machine name and the *slow* variable.
count(downloads3, machineName, slow)

## # A tibble: 10 x 3
##   machineName slow      n
##   <chr>      <chr> <int>
## 1 cs18      No     3662
## 2 cs18      Yes     152
## 3 kermit    No     8717
## 4 kermit    Yes     377
## 5 piglet    No    10734
## 6 piglet    Yes     466
## 7 pluto     No     4963
## 8 pluto     Yes     290
## 9 tweetie   No     6983
## 10 tweetie  Yes     364
```

---

## Sorting data: arrange

The `arrange` function can be used to sort the data according to one or more columns.

Let's sort the data according to download size (ascending order). The first lines of the sorted data set is printed on-screen, but the dataset `downloads3` has *not* been changed.

```
arrange(downloads3, size)

## # A tibble: 36,708 x 6
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18      3   3.73 0.804  0.477 Yes
## 2 piglet    3   1.53 1.96   0.477 Yes
## 3 piglet    3   1.53 1.96   0.477 Yes
## 4 tweetie   3   1.11 2.71   0.477 Yes
## 5 kermit    3   1.12 2.69   0.477 Yes
## 6 pluto     3   8.60 0.349  0.477 Yes
## 7 pluto     3   9.87 0.304  0.477 Yes
## 8 pluto     3   3.78 0.793  0.477 Yes
## 9 pluto     3   4.68 0.641  0.477 Yes
## 10 pluto    3   4.93 0.608  0.477 Yes
```



```
## # ... with 36,698 more rows
```

Two different examples:

```
# According to download size in descending order
arrange(downloads3, desc(size))
```

```
## # A tibble: 36,708 x 6
##   machineName      size    time    speed logSize slow
##   <chr>          <dbl> <dbl>   <dbl>   <dbl> <chr>
## 1 kermit       14518894 1380.   10522.    7.16 No
## 2 piglet       14158123  123.   115169.    7.15 No
## 3 kermit        9523767  171.    55562.    6.98 No
## 4 piglet        9384067   80.0  117309.    6.97 No
## 5 pluto         8674562 1878.    4619.    6.94 No
## 6 kermit        6935603   88.2   78655.    6.84 No
## 7 cs18          6360764   863.    7374.    6.80 No
## 8 piglet        5143062   597.    8611.    6.71 No
## 9 piglet        4812334   215.   22345.    6.68 No
## 10 cs18         4720220 1749.    2700.    6.67 No
## # ... with 36,698 more rows
```

```
# After machine name and then according to download size in descending order
arrange(downloads3, machineName, desc(size))
```

```
## # A tibble: 36,708 x 6
##   machineName      size    time    speed logSize slow
##   <chr>          <dbl> <dbl>   <dbl>   <dbl> <chr>
## 1 cs18          6360764   863.    7374.    6.80 No
## 2 cs18          4720220 1749.    2700.    6.67 No
## 3 cs18          4055821 1275.    3180.    6.61 No
## 4 cs18          3047343   20.9  146038.    6.48 No
## 5 cs18          2952381   318.    9289.    6.47 No
## 6 cs18          2743516 1151.    2383.    6.44 No
## 7 cs18          2691689   0.834 3228695.    6.43 No
## 8 cs18          2613025   18.5  140959.    6.42 No
## 9 cs18          2573336 1335.    1928.    6.41 No
## 10 cs18         1931453   186.   10388.    6.29 No
## # ... with 36,698 more rows
```

---

## Grouping: group\_by

We can group the dataset by one or more categorical variables with `group_by`. The dataset is not changed as such, but - as we will see - grouping can be useful for computation of summary statistics and graphics.

Here we group after machine name (first) *and* the slow variable (second). The only way we can see it at this point is in the second line in the output (`# Groups:`):

```
# Group according to machine
group_by(downloads3, machineName)
```

```
## # A tibble: 36,708 x 6
## # Groups:   machineName [5]
##   machineName      size    time    speed logSize slow
##   <chr>          <dbl> <dbl>   <dbl>   <dbl> <chr>
## 1 cs18          2464 0.493  4998.    3.39 No
```

```
## 2 cs18      7745 0.326 23786.    3.89 No
## 3 cs18      6727 0.314 21444.    3.83 No
## 4 cs18     13049 0.583 22400.    4.12 No
## 5 cs18       356 0.259  1373.    2.55 No
## 6 cs18     15063 0.336 44897.    4.18 No
## 7 cs18      2548 0.285  8945.    3.41 No
## 8 cs18      1932 0.286  6763.    3.29 No
## 9 cs18      7294 0.397 18368.    3.86 No
## 10 cs18     4470 3.41  1311.    3.65 No
## # ... with 36,698 more rows
```

```
# Group according to machine and slow
group_by(downloads3, machineName, slow)
```

```
## # A tibble: 36,708 x 6
## # Groups:   machineName, slow [10]
##   machineName size time speed logSize slow
##   <chr>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1 cs18      2464 0.493 4998.    3.39 No
## 2 cs18      7745 0.326 23786.    3.89 No
## 3 cs18      6727 0.314 21444.    3.83 No
## 4 cs18     13049 0.583 22400.    4.12 No
## 5 cs18       356 0.259  1373.    2.55 No
## 6 cs18     15063 0.336 44897.    4.18 No
## 7 cs18      2548 0.285  8945.    3.41 No
## 8 cs18      1932 0.286  6763.    3.29 No
## 9 cs18      7294 0.397 18368.    3.86 No
## 10 cs18     4470 3.41  1311.    3.65 No
## # ... with 36,698 more rows
```

---

## Summary statistics, revisited: `summarize`

Recall how we could compute summary statistics for a single variable in a dataset, e.g.

```
mean(downloads3$size)
```

```
## [1] 16638.36
```

```
max(downloads3$size)
```

```
## [1] 14518894
```

With `summarize` we can compute summary statistics for a variable for each level of a grouping variable or for each combination of several grouping variables.

First, a bunch of summaries for the size variable for each machine name, where we give explicit names for the new variables:

```
downloads.grp1 <- group_by(downloads3, machineName)
summarize(downloads.grp1,
  avg = mean(size),
  med = median(size),
  stdev = sd(size),
  total = sum(size),
  n = n())
```

```
## # A tibble: 5 x 6
```

```
## machineName avg med stdev total n
## <chr> <dbl> <dbl> <dbl> <dbl> <int>
## 1 cs18 26375. 1990. 208915. 100593281 3814
## 2 kermit 19247. 2466 213985. 175032552 9094
## 3 piglet 14121. 2146. 188340. 158149841 11200
## 4 pluto 13822. 2069 144425. 72605544 5253
## 5 tweetie 14207. 2197 94318. 104379794 7347
```

Second, the same thing but for each combination of machine name and the slow variable:

```
downloads.grp2 <- group_by(downloads3, machineName, slow)
summarize(downloads.grp2,
  avg = mean(size),
  med = median(size),
  stdev = sd(size),
  total = sum(size),
  n = n())
```

```
## # A tibble: 10 x 7
## # Groups: machineName [5]
## machineName slow avg med stdev total n
## <chr> <chr> <dbl> <dbl> <dbl> <dbl> <int>
## 1 cs18 No 27445. 2092. 213140. 100503042 3662
## 2 cs18 Yes 594. 368. 614. 90239 152
## 3 kermit No 20030. 2598 218529. 174602282 8717
## 4 kermit Yes 1141. 541 3049. 430270 377
## 5 piglet No 14687. 2264 192365. 157650747 10734
## 6 piglet Yes 1071. 416. 1934. 499094 466
## 7 pluto No 14564. 2164 148551. 72280790 4963
## 8 pluto Yes 1120. 413 2108. 324754 290
## 9 tweetie No 14894. 2373 96694. 104001733 6983
## 10 tweetie Yes 1039. 471 2603. 378061 364
```

Third, mean and standard deviation for several variables:

```
summarize_at(downloads.grp2, c("time", "size"), list(ave=mean, stdev=sd))
```

```
## # A tibble: 10 x 6
## # Groups: machineName [5]
## machineName slow time_ave size_ave time_stdev size_stdev
## <chr> <chr> <dbl> <dbl> <dbl> <dbl>
## 1 cs18 No 5.17 27445. 57.1 213140.
## 2 cs18 Yes 9.63 594. 17.8 614.
## 3 kermit No 3.41 20030. 25.3 218529.
## 4 kermit Yes 20.7 1141. 47.8 3049.
## 5 piglet No 2.33 14687. 13.8 192365.
## 6 piglet Yes 19.4 1071. 40.2 1934.
## 7 pluto No 3.40 14564. 30.4 148551.
## 8 pluto Yes 21.7 1120. 46.3 2108.
## 9 tweetie No 2.68 14894. 17.3 96694.
## 10 tweetie Yes 17.8 1039. 34.5 2603.
```

The datasets with summaries can be saved as datasets themselves, for example to be used as the basis for certain graphs.

## The pipe operator: %>%

Two or more function calls can be evaluated sequentially using the so-called pipe operator, %>%. Nesting of function calls becomes more readable, and intermediate assignments are avoided.

Let's try it to do a bunch of things in one go, starting with the original dataset:

```
downloads %>%  
  filter(size>0) %>% # Subset of data  
  group_by(machineName) %>% # Grouping  
  summarize(avg = mean(size)) %>% # Compute mean  
  arrange(avg) # Sort after mean
```

```
## # A tibble: 5 x 2  
##   machineName    avg  
##   <chr>        <dbl>  
## 1 pluto        13822.  
## 2 piglet       14121.  
## 3 tweetie      14207.  
## 4 kermit       19247.  
## 5 cs18         26375.
```